

Can Vision-Language-Action Models Adapt to Visual Perturbations Using a Test-Time Trajectory?

Anonymous authors

Abstract—Recently, vision-language-action models (VLAs) have shown strong generalization across diverse robotic tasks. However, when deployed in new environments with unseen visual conditions, their performance often degrades due to visual distribution shifts. Such shifts disrupt the alignment between visual and action representations, leading to distorted action semantics at test-time. Existing approaches primarily address this issue by scaling multi-modal training data or incorporating reasoning mechanisms. However, they require substantial data collection and computational resources, limiting their practicality for various real-world deployments. To address this challenge, we propose TTA-VLA, a simple yet effective test-time adaptation framework for VLAs using only a test-time trajectory. Our method leverages environment-specific observations from this trajectory as self-supervision signals for adaptation. Based on these observations, we extract visual and action embedding transitions from a pre-trained VLA and align them to reduce visual-action temporal inconsistencies induced by unseen visual shifts. To preserve task semantics in the representations under visual distribution shifts, we introduce a language-conditioned reference that grounds visual-action transition features in the task instruction, providing an environment-invariant anchor. Across various visual perturbation scenarios on benchmark and real-world experiments, our method achieves an average success rate improvement of 24% over the pre-trained policies.

I. INTRODUCTION

Vision-language-action models (VLAs) have shown remarkable progress in enabling robots to interpret natural language instructions from humans, understand complex visual scenes, and execute manipulation behaviors [1], [2], [3], [4], [5]. They jointly learn from large-scale robotic multi-modal dataset [6], [7], [8], including images, language instructions, robot sensory data, and robot demonstrations. As a result, VLAs can integrate the capability of perception, reasoning, and action generation into an end-to-end policy, highlighting their potential as a scalable and flexible foundation for generalist robot learning [9], [10].

Despite their success, VLAs still suffer from critical limitations when deployed in new environments [3], [4], [12]. In real-world, visual perturbations such as background texture, lighting condition, and camera viewpoints can readily arise and deviate from the training distribution, as illustrated in Fig. 1. [11] shows that existing VLAs are highly sensitive to these perturbations, with average success rates dropping from around 95% to below 30%. Consistently, we observe in our environment that even minor appearance shifts can severely degrade policy performance, yielding an average 70% drop across the six perturbation categories, as in Fig. 1h.

To address these problems, most recent works have mainly focused on advancing the training recipe. A common direc-




Fig. 1: Examples of various visual perturbation types in LIBERO-Plus [11], along with their performance on Open-VLA [4]. Original performance without visual perturbation, exemplified by (a), is plotted as a dashed line in (h).

tion is to increase the scale and diversity of demonstrations collected across varied scenes, objects, and sensing conditions [3], [4], [12], [13]. A broad range of observations and manipulation behaviors helps the model to learn perception-action relationships that generalize more consistently in novel environments. Another direction is to enhance the reasoning capability of the underlying vision-language backbone [5], [14], [15], [16]. Reasoning mechanisms enable VLAs to generate grounded intermediate representations, such as object relations and sub-task decomposition, allowing them to leverage the semantic and compositional capabilities of pre-trained large-language models and visual-language models.

While these two directions offer meaningful benefits, they aim to improve the capability of VLAs through extensive training, resulting in substantial additional data collection and computational resources. Furthermore, these advances typically strengthen visual understanding, rather than explicitly addressing distribution shifts at deployment. Consequently, real-world environments vary widely and may include conditions not covered during training, which can lead to inconsistent action behaviors under the resulting distribution shifts, as illustrated in Fig. 2a and Fig. 2b. These considerations motivate investigating complementary approaches that enable models to adapt using information available at test time, allowing more reliable operation under unseen conditions.




Fig. 2: Comparison of VLAs under in-distribution and out-of-distribution. Blue denotes the in-distribution training domain, green indicates the out-of-distribution domain at test time, and orange represents inconsistent action outputs induced by distribution shifts. (a) In an in-distribution setting, observations align with the training distribution of the VLA, resulting in consistent visual–action representations. (b) Under an out-of-distribution environment, visually shifted observations lead to mismatches between visual inputs and action representations in the pre-trained VLA. (c) Our method effectively adapts the pre-trained VLA using a test-time trajectory in the shifted environment, mitigating visual–action inconsistencies.

To adapt the pre-trained model during deployment without extensive training, test-time adaptation (TTA) provides a natural solution across various domains [17], [18], [19]. It adjusts the model parameters using only the inputs available at test time, typically optimizing self-supervised objectives that encourage confident predictions with entropy minimization, consistency regularization, or pseudo-labeling [20], [21], [22], [23]. Despite its potential benefits, applying test-time adaptation to VLAs poses additional challenges due to their multi-modal characteristics. In particular, VLAs depend on tight alignment across vision, language, and action modalities. As a result, even a modest visual distribution shift can easily disrupt this alignment and compromise downstream action semantics. However, prior TTA approaches in robotics do not consider the alignment of multi-modal representations, only relying on feedback or entropy of predictions as output-level supervision signals [17], [24]. Furthermore, addressing such cross-modal misalignment at test time requires supervision signals at the representation-level that reflect whether actions remain semantically grounded in both the instruction and the scene. Yet, although robots can obtain rich raw sensory observations at deployment, it remains challenging to derive test-time supervision that explicitly captures cross-modal semantic grounding among vision, language, and action, since existing TTA objectives are largely agnostic to these semantics.

To overcome these challenges, we introduce a **Test-time Trajectory-based Adaptation framework for VLAs**, named **TTA-VLA**. It exploits visual observations collected from a trajectory in the new deployment settings to explicitly align visual and action representations, as shown in Fig. 2c. The key insight behind TTA-VLA is that during training, VLAs learn relationships between input observations and output actions, yielding transition patterns in which changes in visual observations induce corresponding action transitions over time. We leverage this temporal consistency between visual and action transitions in their latent representations as a self-supervised signal for test-time adaptation. Under visual domain shifts, however, the pre-trained VLA may fail to preserve its original observation–action relationship, leading to reduced visual–action temporal consistency such

that visual changes no longer induce appropriate action transitions. To address this, we extract transition representations from consecutive visual and action embeddings along a test-time trajectory as environment-specific information. Then, we align these transition representations with the task instruction, which serves as a language-conditioned semantic reference, to reduce visual–action temporal inconsistencies induced by unseen visual shifts. This alignment grounds the self-supervised adaptation signal in vision, language, and action semantics, encouraging both visual and action transitions to remain semantically consistent with the intended task. In this way, the proposed method provides a simple yet effective test-time adaptation that restores disrupted visual–action alignment while maintaining semantic stability under distribution shifts, without requiring supervision from experts or computational overhead. This formulation is not tailored to a specific VLA architecture, but instead addresses test-time adaptation as a general challenge arising from cross-modal misalignment between visual and action representations.

We validate the effectiveness of our test-time adaptation framework under a wide range of visual distribution shifts on two representative VLA paradigms: the autoregressive-based OpenVLA [4] and the flow-matching-based π_0 [3]. For simulation experiments on the LIBERO-Plus [11] with six visual perturbation types, we observe consistent performance improvements over the pre-trained policy across both VLA architectures with an average gain of 14%. To assess the real-world practicality of our approach, we evaluate it under four visual perturbation deployment scenarios, achieving an average success rate improvement of 33%.

II. RELATED WORKS

A. Vision-language-action models

Built upon the strong compositionality and semantic grounding capabilities of pre-trained vision-language models (VLMs) [25], [26], vision-language-action models (VLAs) achieve powerful generalist robot control by taking sensory observations and language instructions as input and producing low-level actions as output [9], [10]. Recent VLAs adopt diverse action-generation architectures, which can be broadly categorized into three groups: (i) autoregressive poli-

cies that extend VLMs to sequential action generation [1], [4], [27], [28], (ii) diffusion-based methods, which produce continuous actions through iterative denoising [2], [3], [29], [30], and (iii) CLIP-based approach using similarity between visual context and language-based motions [12]. To further strengthen zero-shot generalization, prior works primarily expand the coverage and diversity of multi-modal robotic training datasets, enabling VLAs to generalize across various situations [3], [4], [12], [13]. More recently, several methods augment VLAs with high-level reasoning mechanisms, such as subtask-conditioned action generation [5], [14], [15], [16].

However, they primarily focus on improving robustness of VLAs through large-scale training with almost 1 million demonstrations [7] or additional reasoning data [16], requiring extensive computational resources. Furthermore, real-world deployments may expose models to conditions beyond the training distribution due to its diversity, inducing distribution shifts that can degrade performance at test-time. To address this, our work focuses on test-time adaptation, explicitly re-aligning disrupted relationships of observation and action representations through a test-time trajectory.

B. Test-time adaptation

Test-time adaptation (TTA) aims to mitigate distribution shift by adjusting the model parameters at test-time using only unlabeled test-time data, typically by optimizing self-supervised objectives that encourage consistent and confident predictions, thereby avoiding costly retraining or additional supervision [31], [32]. TTA has been widely investigated for maintaining feature consistency under distribution shift with entropy minimization [19], [22] or batch-norm statistics adaptation [23], [33]. For vision-language models, prior works explore optimizing text prompts using augmented images [18], [20] or employing self-supervised proxy tasks [21], [34]. In robotics, only a few studies have explored TTA, primarily in visual-language navigation, using entropy minimization of action predictions [24] or binary feedback-based reinforcement learning [17], which exploit output-level signals rather than the internal latent representations of the model. However, TTA for vision-language-action models remains unexplored, despite their susceptibility to visual distribution shifts.

In this work, we introduce test-time adaptation of VLAs by leveraging a test-time trajectory to self-supervise the alignment between visual and action representations disrupted by visual distribution shifts. To stabilize this alignment with task semantics, we incorporate the task instruction as a language-conditioned reference, encouraging visual and action features to remain semantically consistent at the feature-level. In VLAs, test-time scaling has been studied to improve action generation at test-time by producing multiple candidates followed by post-hoc selection [35], which is orthogonal to our focus because it only changes the inference procedure rather than adapting model representations disrupted by visual distribution shifts.

III. PROPOSED METHOD

A. Problem statement: test-time adaptation of VLAs

Our goal is to enable vision-language-action models (VLAs) to maintain reliable action generation under visual distribution shifts by adapting to new environments using only a test-time trajectory. Typically, VLAs, initialized from pre-trained VLMs, enable general-purpose robot control by taking sensory observations and a task instruction as input and producing low-level actions as output. VLAs are trained on a large-scale dataset of expert demonstrations $\mathcal{D}_{train} = \{(o_t, i_t, a_t)\}_{t=1}^N$, where o_t , i_t , and a_t at time step t denote sensory observation, natural-language task instruction, and its corresponding expert action, respectively. Note that the observation is an RGB image, following standard VLA setup [4]. The image is first encoded into a visual feature through the visual encoder $\mathcal{F}_v$ as $\bar{z}_t = \mathcal{F}_v(o_t)$ and then transformed by a visual projector $\mathcal{V}$ to make it compatible with the language model, $z_t = \mathcal{V}(\bar{z}_t)$. The task instruction is also embedded using a text encoder $\mathcal{F}_l$ as $e_t = \mathcal{F}_l(i_t)$. The projected visual feature z_t and language feature e_t are concatenated and fed into the language model, which produces a latent representation u_t that is mapped to the robotic action space. In this way, the VLA model parameterized with θ learns the observation-action relationships by minimizing the supervised objective with the action prediction loss $\mathcal{L}_\theta$:

$$\min_{\theta} \mathbb{E}_{(o_t, i_t, a_t) \sim \mathcal{D}_{train}} [\mathcal{L}_\theta(a_t | o_t, i_t)]. \quad (1)$$

When VLAs are deployed in a new environment, they encounter new sensory observations in $\mathcal{D}_{new}$ that differ from the distribution of training data $\mathcal{D}_{train}$ due to the changes in background, lighting, viewpoint, or object appearance. In the context of VLAs, such visual domain shifts do not merely distort visual features as typically in vision tasks, they fundamentally disrupt the tri-modal alignment between vision, language, and action. Specifically, the visual embedding z_t becomes misaligned from the joint embedding space where language instructions e_t and action sequences a_t are anchored, leading to action predictions that are no longer semantically grounded. Therefore, we define the primary challenge of TTA for VLAs is to adapt visual representations to the new environment while preserving the cross-modal semantic consistency.

At test-time, however, VLAs no longer have access to expert actions, ground-truth labels, or additional demonstrations that could be used to recover this misalignment. The only available signal is incoming sensory observations from the new environment, which may significantly differ from the training distribution. Consequently, we formulate the TTA problem as the following unsupervised optimization:

$$\theta = \operatorname{argmin}_{\theta_o} \mathbb{E}_{(o_t, i_t) \sim \mathcal{D}_{new}} [\mathcal{L}_{\theta_o}^{\text{TTA}}(o_t, i_t, \theta_o)], \quad (2)$$

where $\mathcal{L}_{\theta_o}^{\text{TTA}}$ denotes an objective function for test-time adaptation, and θ_o represents the pre-trained model parameters. This adjusts the pre-trained parameters by leveraging its behavior when exposed to observations from a new distribution.




Fig. 3: Overview. The robot collects observations by executing a trajectory in the new environment. Each observation is converted into visual and action features, z and u , through the pre-trained VLA. Our proposal aligns their transitions, Δz_{t+1} and Δu_{t+1} , while anchoring this alignment to the task-instruction embedding e_t to ensure that both transitions remain semantically consistent with the task i_t .

B. TTA-VLA: Test-time Trajectory-based Adaptation for VLAs

To address the above-mentioned challenges in tri-modal misalignment, we propose TTA-VLA, which leverages the temporal structure of observation–action relationships conditioned on language as a self-supervised signal. Specifically, we obtain this signal from a test-time robot trajectory collected in the new environment by arbitrarily executing the robot, from which we capture how the new environment affects the pre-trained behaviors of the VLA, as depicted in Fig. 3. Formally, this trajectory can be written as a sequence of observation–action pairs, $(o_t, a_t, o_{t+1}, a_{t+1}, \dots)$. From this trajectory, we encode the observations to extract visual features from the visual projector, z_t and z_{t+1} , and action features from the language model, u_t and u_{t+1} . We then compute their temporal transitions as follows:

$$\Delta z_{t+1} = z_{t+1} - z_t \quad \text{and} \quad \Delta u_{t+1} = u_{t+1} - u_t. \quad (3)$$

Our key insight for feature alignment is that VLAs implicitly learn a temporal correspondence between changes in observations and the action representations during training. However, visual domain shifts disrupt this correspondence, motivating the restoration of consistency between the visual transition Δz_{t+1} and the action transition Δu_{t+1} for maintaining reliable behavior of VLAs in the new environment.

To this end, we introduce a transition-based adaptation loss as $\mathcal{L}_{\text{align}}(\Delta z_{t+1}, \Delta u_{t+1})$ to reduce the temporal inconsistency between visual and action feature transitions. We investigate two geometric alignment strategies, angular-based and projection-based, that encourage directional consistency between transitions across modality spaces. The angular-based alignment focuses on matching the orientation of feature transitions by minimizing their cosine distance, inspired by [36]:

$$1 - \frac{\Delta z_{t+1}^\top \Delta u_{t+1}}{|\Delta z_{t+1}|_2 |\Delta u_{t+1}|_2}. \quad (4)$$

The projection-based variant projects the action feature transition onto the visual feature transition and penalizes the residual, encouraging action transitions to follow visual transitions, inspired by [37], with a constant for numerical stability ϵ :

$$\left\| \Delta u_{t+1} - \frac{\Delta u_{t+1}^\top \Delta z_{t+1}}{\|\Delta z_{t+1}\|_2^2 + \epsilon} \Delta z_{t+1} \right\|_2^2. \quad (5)$$

Their corresponding analyses are provided in Section IV-A.2.

To further strengthen the language semantic grounding of the adapted representation, we introduce a language-conditioned constraint, which uses the natural-language task description as an environment-invariant reference point. Since the task instruction defines the intended goal and remains unchanged across visual domains, its embedding provides a stable language semantic anchor that guides test-time adaptation while mitigating semantic drift. We incorporate the task instruction as a language semantic anchor by re-centering the visual and action transitions around the text embedding. This gives the following transitions:

$$\Delta z_{t+1}^{\text{anch}} = \Delta z_{t+1} - e_t \quad \text{and} \quad \Delta u_{t+1}^{\text{anch}} = \Delta u_{t+1} - e_t, \quad (6)$$

where e_t denotes the text embedding of the task instruction i_t extracted from the text encoder $\mathcal{F}_l$. Then, we rewrite the objective function as follows:

$$\mathcal{L}_{\theta_o}^{\text{TTA}} = \mathcal{L}_{\text{align}}(\Delta z_{t+1}^{\text{anch}}, \Delta u_{t+1}^{\text{anch}}). \quad (7)$$

By grounding the transition dynamics in the task semantics, the proposal not only stabilizes adaptation but also prevents semantic drift, enabling more reliable behavior under visual domain shifts. During adaptation, only the parameters of the visual projector $\mathcal{V}$ are updated, while all other components of the VLA policy remain frozen, following [38], [39].

IV. EXPERIMENTS

A. Benchmark experiments

1) *Setup: Dataset.* We validate our proposal on LIBERO-Plus [11], which extends LIBERO [40] with systematic and realistic visual perturbations for evaluating test-time adaptation under visual distribution shifts. We assess performance across six axes of visual domain shifts, including background texture, camera viewpoint, lighting condition, robot initial state, sensor noise, and object layout, with examples shown in Fig. 1. These perturbations expand the original set of 10 tasks to approximately 353 tasks on average across the object, spatial, and goal variants of the benchmark.

Implementation details. We evaluate our method on two representative VLAs with different action-generation designs: autoregressive-based OpenVLA [4] and flow-matching-based π_0 [3], both of which are widely used, open-source, and published in peer-reviewed venues. We compare our method against the corresponding pretrained VLA without adaptation under two conditions: (i) in-distribution settings without visual perturbations and (ii) out-of-distribution settings with visual perturbations, to assess both the robustness of pretrained VLAs to visual distribution shifts and the

TABLE I: Comparison results of task success rates on LIBERO-Plus across six visual perturbation scenarios.

Method	Background texture	Camera viewpoint	Light condition	Robot initial state	Sensor noise	Object layout	Avg.
OpenVLA							
In-distribution	90.0%						90.0%
Out-of-distribution	46.4%	1.2%	11.4%	3.7%	5.9%	46.4%	19.2%
TTA-VLA (ours)	66.3%	2.7%	26.6%	11.3%	10.7%	62.9%	30.1%
π_0							
In-distribution	90.0%						90.0%
Out-of-distribution	79.8%	15.1%	82.5%	5.8%	15.6%	72.6%	45.2%
TTA-VLA (ours)	83.1%	46.2%	88.5%	17.7%	61.0%	75.9%	62.1%




Fig. 4: Task success rates according to task difficulty under visual perturbations, where difficulty levels in x-axis follow the original LIBERO-Plus. Orange and blue lines denote the original policy and our TTA-VLA policy, respectively.

effectiveness of our adaptation method. We use the publicly available pre-trained policy for each backbone released by the authors and follow the learning hyperparameters reported in the corresponding papers during adaptation.

2) *Results: Main results.* We summarize the comparison results on six visual perturbation scenarios in TABLE I, reporting success rates averaged over object, spatial, and goal variants of LIBERO-Plus. While both OpenVLA and π_0 achieve high performance in the in-distribution setting as 90.0% (first rows) in our environment, their success rates decrease substantially under visual perturbations (second rows), with average drops of 70.8% and 44.8%, respectively, highlighting the sensitivity of pre-trained VLAs to visual distribution shifts. Their performance drop is larger for OpenVLA than for π_0 , likely due to its reliance on a single-camera view. In contrast, background texture and object layout perturbations exhibit the smallest drops among the six types, due to broader exposure to appearance and layout variations during training. The proposed method, TTA-VLA, achieves the best performance across all settings, with improvements of up to 45.4% over the method without adaptation under perturbations. Notably, the benefits are most pronounced under the most challenging conditions, including camera viewpoint, robot initial state, and sensor noise, where the proposed method achieves substantial additional improvements of 16.3%, 9.8%, and 25.1%, respectively, when averaged across the two backbones. Overall, these results demonstrate that our transition-based adaptation incorporating language-conditioned constraints helps align visual and action representations with the intended task semantics, encouraging both modalities to remain semantically consistent throughout test-time execution under severe visual distribution shifts.

Performance according to task difficulty. We analyze our method with respect to task difficulty, as defined in LIBERO-

Plus. In this benchmark, task instances are grouped into five difficulty levels based on how many evaluated VLAs can complete them, ranging from L1, where all models succeed, to L5, where none succeed. Fig. 4 reports task success rates across difficulty levels for each perturbation type, averaged over two backbones, comparing the original and adapted policies. Performance decreases from L1 to L5 across all perturbations for both policies. Nevertheless, our method consistently outperforms the original policy across all difficulty levels, with average improvements of 16.8% and 10.8% at L1 and L5, respectively, across the six perturbations. Notably, for background texture, camera viewpoint, and object layout, the performance gap widens at higher difficulty levels (L3-L5), indicating effectiveness on challenging executions that are inherently difficult for original policies. These results suggest that TTA-VLA not only preserves performance on easy tasks but also enhances robustness on difficult task instances under visual distribution shifts by leveraging test-time trajectories to mitigate compounded challenges from both task difficulty and visual perturbations.

Knowledge preservation and improvement. To assess the stability and plasticity trade-off of our test-time adaptation framework, we analyze the extent to which successful behaviors of the pre-trained policy are preserved and failed ones are improved under visual distribution shifts. Fig. 5 reports the corresponding results across two models. *Preservation* denotes the percentage of originally successful tasks that remain successful after adaptation, while *Improvement* denotes the percentage of originally failed tasks that become successful after adaptation. Our approach preserves most tasks successfully completed by the original policy, achieving an average preservation of 85%. At the same time, our proposal substantially improves previously failed cases, yielding an average improvement of 26%. Interestingly, background




Fig. 5: Knowledge preservation and improvement under various visual perturbations compared to the original policy.

TABLE II: Effect of the text-anchored alignment. ‘Text’ denotes the language-conditioned constraint.

Method	OpenVLA	π_0	Avg.
w/o Text	29.1%	60.5%	44.8%
w/ Text	30.1%	62.1%	46.1%

texture, sensor noise, and object layout perturbations exhibit the highest levels of both preservation and improvement among the six perturbation types. These results indicate that our method effectively balances maintaining existing successful behaviors and recovering failures, demonstrating robust adaptation without sacrificing prior performance under diverse visual distribution shifts.

Effect of text-anchored alignment. To analyze the contribution of text-anchored alignment during adaptation, we compare the full TTA-VLA objective in Eq. (7) with an unanchored variant that directly aligns visual and action transitions without the language-conditioned constraint. As shown in TABLE II, the method with the language-conditioned constraint consistently yields an average gain of 1.3% across two backbones. These results indicate that transition-based alignment provides the primary self-supervised signal for adaptation by enforcing temporal consistency between visual observations and actions, while the language-conditioned reference further stabilizes the adapted representations by preserving task semantics.

Variations of alignment strategies. We further compare projection-based and angular-based feature alignment in $\mathcal{L}_{\text{align}}$ using OpenVLA. More details of alignment approaches in Section III-B. As shown in TABLE III, the projection-based strategy consistently outperforms the angular-based one across all settings, with average gains of 3.5%. This suggests that explicitly suppressing deviated transition components provides more stable alignment than only constraining angular direction. However, the gap between the two strategies is relatively small compared with the overall improvement over the original policy.

Visualizations of rollouts. Fig. 6 presents comparisons of rollouts between the original policy and TTA-VLA under light condition and sensor noise perturbations. Under light condition, the original policy fails to execute precise manipulation, exhibiting inaccurate grasp poses that prevent stable pickup of the target bowl. In contrast, our method

TABLE III: Performance improvements with projection-based alignment compared to angular-based alignment.

Object	Spatial	Goal	Avg.
+ 2.4	+ 3.4	+ 4.2	+ 3.5




Fig. 6: Visualizations of rollouts with the original policy (top row) and TTA-VLA (bottom row) on LIBERO-Plus under (a) light condition and (b) sensor noise perturbations. Key failure scenes are highlighted with red circles, while successful executions are highlighted with blue boxes.

successfully adapts the test-time visual perturbation, producing accurate grasp configurations and completing the task reliably despite the challenging illumination changes. Similarly, under sensor noise, the original policy struggles to localize the target object, leading to incorrect reaching behaviors and task failure. However, TTA-VLA identifies the target object and executes the pick-and-place sequence correctly, even under severe observation corruption. More rollout visualizations under other perturbation types are provided in the multimedia materials.

B. Real-world experiments

1) Setup: Tasks. To evaluate robustness while isolating the effects of visual perturbations and adaptation from other factors affecting task performance, we use four fundamental tabletop manipulation tasks representing core manipulation capabilities: pick-and-place, slide, push, and pull. The first task, *Pick up the red block and place it on the plate* (Pick-and-place), requires the robot to localize the target object, grasp it stably, and place it at the designated location. The second task, *slide the green truck to the red car* (Slide), expects the robot to maintain contact with the object and slide it along the surface toward the target. The third task, *push the blue button* (Push), requires the robot to position its end

TABLE IV: Comparison results of task success rates (%) on real-world experiments with four visual perturbation scenarios.

Task	In-distribution	Background		Camera		Robot		Object		Avg.	
		Org.	Ours	Org.	Ours	Org.	Ours	Org.	Ours	Org.	Ours
Pick-and-place	100	50	90	15	45	35	65	70	100	39	75
Slide	100	70	100	25	60	15	60	70	100	45	80
Push	90	90	95	25	70	50	80	90	100	64	86
Pull	90	70	90	0	40	40	90	80	95	48	79
Avg.	95	70	95	22	67	31	74	78	97	50	83




Fig. 7: Real-world rollouts between the original policy (first row) and our method (second row). Key failure scenes are highlighted with red circles, while successful executions are highlighted with blue boxes.

effector and apply a directional pushing motion. The fourth task, *pull out the tissue* (Pull), expects the robot to grasp the tissue and pull it out smoothly without losing contact.

Implementation details. We conduct real-world experiments using a 7-DoF robotic arm Franka Research 3 (FR3). For the evaluation, we employ OpenVLA using visual input from a frontal third-person camera and the language task instruction. We design four types of visual perturbations, including background texture (Background), camera view-point (Camera), robot initial state (Robot), and object layout (Object). For Background, the in-distribution setting uses a black background, while five additional backgrounds with different colors are used for evaluation. For Camera, the in-distribution setting uses a frontal third-person view, and four additional camera viewpoints are evaluated. For Robot, the in-distribution setting corresponds to the home pose, while five alternative initial poses are used. For Object, we construct five evaluation settings by introducing new surrounding objects or by changing the positions of existing surrounding objects relative to the in-distribution configuration.

2) Results: Quantitative results. TABLE IV summarizes the overall success rates of the original policy and TTA-VLA across four real-world scenarios. The original policy is highly sensitive to visual distribution shifts, especially under Camera and Robot perturbations, with average drops of 73% and 64%, respectively. In contrast, TTA-VLA consistently improves performance across all tasks and perturbation types, achieving an average gain of 33% over the pre-trained policy. The gains are particularly large under Camera and Robot perturbations, reaching 45% and 43%, while remaining consistent under Background and Object shifts with gains of 25% and 19%. These improvements suggest that TTA-VLA generalizes beyond a single skill by restoring visual-action inconsistencies through transition-

based alignment with language-conditioned supervision.

Visualizations of rollouts. Fig. 7 shows representative real-world rollouts of the original policy (top) and TTA-VLA (bottom). The original policy exhibits unstable behaviors, such as mislocalizing the target object under background shifts in Slide and being distracted by changes in surrounding objects in Pick-and-place, which leads to missing the target or interacting with the wrong region. In contrast, the proposed method maintains stable, target-centric motion and completes the intended manipulation by restoring visual-action inconsistencies induced by visual distribution shifts at the representation-level. These results support that transition-based alignment remains effective under real-world distribution shifts by reducing visual-action inconsistencies. More rollouts are provided in the multimedia materials.

V. CONCLUSION

We have addressed the challenge of deploying vision-language-action models in environments with unseen visual conditions. To this end, we have introduced TTA-VLA, a test-time trajectory-based adaptation framework that leverages visual-action transition consistency conditioned on language as a self-supervised signal, enabling VLAs to adapt using a small number of test-time observations. To further preserve semantic stability, it anchors transition alignment to the natural-language task instruction, providing an environment-invariant semantic reference that mitigates representational drift. Through extensive experiments on the benchmark and real-world scenarios, we demonstrate that our method significantly improves success rates under diverse visual distribution shifts. These results highlight that test-time adaptation, driven by simple yet effective self-supervised signals, is a promising direction for enhancing the robustness of VLAs in reliable deployment.

REFERENCES

- [1] S. Belkhale, T. Ding, T. Xiao, P. Sermanet, Q. Vuong, J. Tompson, Y. Chebotar, D. Dwibedi, and D. Sadigh, “Rt-h: Action hierarchies using language,” *arXiv preprint arXiv:2403.01823*, 2024.
- [2] J. Bjorck, F. Castañeda, N. Cherniadev, X. Da, R. Ding, L. Fan, Y. Fang, D. Fox, F. Hu, S. Huang, *et al.*, “Gr00t n1: An open foundation model for generalist humanoid robots,” *arXiv preprint arXiv:2503.14734*, 2025.
- [3] K. Black, N. Brown, D. Driess, A. Esmail, M. Equi, C. Finn, N. Fusai, L. Groom, K. Hausman, B. Ichter, *et al.*, “ π_0 : A vision-language-action flow model for general robot control,” *arXiv preprint arXiv:2410.24164*, 2024.
- [4] M. J. Kim, K. Pertsch, S. Karamcheti, T. Xiao, A. Balakrishna, S. Nair, R. Rafailov, E. Foster, G. Lam, P. Sanketi, *et al.*, “Openvla: An open-source vision-language-action model,” *arXiv preprint arXiv:2406.09246*, 2024.
- [5] G. R. Team, S. Abeyruwan, J. Ainslie, J.-B. Alayrac, M. G. Arenas, T. Armstrong, A. Balakrishna, R. Baruch, M. Bauza, M. Blokzijl, *et al.*, “Gemini robotics: Bringing ai into the physical world,” *arXiv preprint arXiv:2503.20020*, 2025.
- [6] A. Khazatsky, K. Pertsch, S. Nair, A. Balakrishna, S. Dasari, S. Karamcheti, S. Nasiriany, M. K. Srirama, L. Y. Chen, K. Ellis, *et al.*, “Droid: A large-scale in-the-wild robot manipulation dataset,” *arXiv preprint arXiv:2403.12945*, 2024.
- [7] A. O’Neill, A. Rehman, A. Maddukuri, A. Gupta, A. Padalkar, A. Lee, A. Pooley, A. Gupta, A. Mandlekar, A. Jain, *et al.*, “Open x-embodiment: Robotic learning datasets and rt-x models: Open x-embodiment collaboration 0,” in *2024 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2024, pp. 6892–6903.
- [8] H. R. Walke, K. Black, T. Z. Zhao, Q. Vuong, C. Zheng, P. Hansen-Estruch, A. W. He, V. Myers, M. J. Kim, M. Du, *et al.*, “Bridgedata v2: A dataset for robot learning at scale,” in *Conference on Robot Learning*. PMLR, 2023, pp. 1723–1736.
- [9] Y. Ma, Z. Song, Y. Zhuang, J. Hao, and I. King, “A survey on vision-language-action models for embodied ai,” *arXiv preprint arXiv:2405.14093*, 2024.
- [10] R. Sapkota, Y. Cao, K. I. Roumeliotis, and M. Karkee, “Vision-language-action models: Concepts, progress, applications and challenges,” *arXiv preprint arXiv:2505.04769*, 2025.
- [11] S. Fei, S. Wang, J. Shi, Z. Dai, J. Cai, P. Qian, L. Ji, X. He, S. Zhang, Z. Fei, *et al.*, “Libero-plus: In-depth robustness analysis of vision-language-action models,” *arXiv preprint arXiv:2510.13626*, 2025.
- [12] G.-C. Kang, J. Kim, K. Shim, J. K. Lee, and B.-T. Zhang, “Clip-rt: Learning language-conditioned robotic policies from natural language supervision,” *arXiv preprint arXiv:2411.00508*, 2024.
- [13] Y. Hu, F. Lin, P. Sheng, C. Wen, J. You, and Y. Gao, “Data scaling laws in imitation learning for robotic manipulation,” *arXiv preprint arXiv:2410.18647*, 2024.
- [14] Y. Ji, H. Tan, J. Shi, X. Hao, Y. Zhang, H. Zhang, P. Wang, M. Zhao, Y. Mu, P. An, *et al.*, “Robobrain: A unified brain model for robotic manipulation from abstract to concrete,” in *Proceedings of the Computer Vision and Pattern Recognition Conference*, 2025, pp. 1724–1734.
- [15] Y. Yang, J. Sun, S. Kou, Y. Wang, and Z. Deng, “Lohovla: A unified vision-language-action model for long-horizon embodied tasks,” *arXiv preprint arXiv:2506.00411*, 2025.
- [16] M. Zawalski, W. Chen, K. Pertsch, O. Mees, C. Finn, and S. Levine, “Robotic control via embodied chain-of-thought reasoning,” *arXiv preprint arXiv:2407.08693*, 2024.
- [17] S. Kim, G. Oh, H. Ko, D. Ji, D. Lee, B.-J. Lee, S. Jang, and S. Kim, “Test-time adaptation for online vision-language navigation with feedback-based reinforcement learning,” in *Forty-second International Conference on Machine Learning*.
- [18] M. Shu, W. Nie, D.-A. Huang, Z. Yu, T. Goldstein, A. Anandkumar, and C. Xiao, “Test-time prompt tuning for zero-shot generalization in vision-language models,” *Advances in Neural Information Processing Systems*, vol. 35, pp. 14 274–14 289, 2022.
- [19] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, and T. Darrell, “Tent: Fully test-time adaptation by entropy minimization,” *arXiv preprint arXiv:2006.10726*, 2020.
- [20] C.-M. Feng, K. Yu, Y. Liu, S. Khan, and W. Zuo, “Diverse data augmentation with diffusions for effective test-time prompt tuning,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023, pp. 2704–2714.
- [21] Y. Liu, P. Kothari, B. Van Delft, B. Bellot-Gurlet, T. Mordan, and A. Alahi, “Ttt++: When does self-supervised test-time training fail or thrive?” *Advances in Neural Information Processing Systems*, vol. 34, pp. 21 808–21 820, 2021.
- [22] C. K. Mummadi, R. Huttmacher, K. Rambach, E. Levinkov, T. Brox, and J. H. Metzen, “Test-time adaptation to distribution shift by confidence maximization and input transformation,” *arXiv preprint arXiv:2106.14999*, 2021.
- [23] Z. Nado, S. Padhy, D. Sculley, A. D’Amour, B. Lakshminarayanan, and J. Snoek, “Evaluating prediction-time batch normalization for robustness under covariate shift,” *arXiv preprint arXiv:2006.10963*, 2020.
- [24] J. Gao, X. Yao, and C. Xu, “Fast-slow test-time adaptation for online vision-and-language navigation,” *arXiv preprint arXiv:2311.13209*, 2023.
- [25] L. Beyer, A. Steiner, A. S. Pinto, A. Kolesnikov, X. Wang, D. Salz, M. Neumann, I. Alabdulmohsin, M. Tschannen, E. Bugliarello, *et al.*, “Paligemma: A versatile 3b vlm for transfer,” *arXiv preprint arXiv:2407.07726*, 2024.
- [26] S. Karamcheti, S. Nair, A. Balakrishna, P. Liang, T. Kollar, and D. Sadigh, “Prismatic vlms: Investigating the design space of visually-conditioned language models,” in *Forty-first International Conference on Machine Learning*, 2024.
- [27] K. Pertsch, K. Stachowicz, B. Ichter, D. Driess, S. Nair, Q. Vuong, O. Mees, C. Finn, and S. Levine, “Fast: Efficient action tokenization for vision-language-action models,” *arXiv preprint arXiv:2501.09747*, 2025.
- [28] B. Zitkovich, T. Yu, S. Xu, P. Xu, T. Xiao, F. Xia, J. Wu, P. Wohlhart, S. Welker, A. Wahid, *et al.*, “Rt-2: Vision-language-action models transfer web knowledge to robotic control,” in *Conference on Robot Learning*. PMLR, 2023, pp. 2165–2183.
- [29] S. Liu, L. Wu, B. Li, H. Tan, H. Chen, Z. Wang, K. Xu, H. Su, and J. Zhu, “Rdt-1b: a diffusion foundation model for bimanual manipulation,” *arXiv preprint arXiv:2410.07864*, 2024.
- [30] J. Wen, Y. Zhu, J. Li, M. Zhu, Z. Tang, K. Wu, Z. Xu, N. Liu, R. Cheng, C. Shen, *et al.*, “Tinyvla: Towards fast, data-efficient vision-language-action models for robotic manipulation,” *IEEE Robotics and Automation Letters*, 2025.
- [31] J. Liang, R. He, and T. Tan, “A comprehensive survey on test-time adaptation under distribution shifts,” *International Journal of Computer Vision*, vol. 133, no. 1, pp. 31–64, 2025.
- [32] Z. Xiao and C. G. Snoek, “Beyond model adaptation at test time: A survey,” *arXiv preprint arXiv:2411.03687*, 2024.
- [33] S. Schneider, E. Rusak, L. Eck, O. Bringmann, W. Brendel, and M. Bethge, “Improving robustness against common corruptions by covariate shift adaptation,” *Advances in neural information processing systems*, vol. 33, pp. 11 539–11 551, 2020.
- [34] Y. Sun, X. Wang, Z. Liu, J. Miller, A. Efros, and M. Hardt, “Test-time training with self-supervision for generalization under distribution shifts,” in *International conference on machine learning*. PMLR, 2020, pp. 9229–9248.
- [35] J. Kwok, C. Agia, R. Sinha, M. Foutter, S. Li, I. Stoica, A. Mirhoseini, and M. Pavone, “Robomonkey: Scaling test-time sampling and verification for vision-language-action models,” in *9th Annual Conference on Robot Learning*.
- [36] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, *et al.*, “Learning transferable visual models from natural language supervision,” in *International conference on machine learning*. PMLR, 2021, pp. 8748–8763.
- [37] B. Fernando, A. Habrard, M. Sebban, and T. Tuytelaars, “Unsupervised visual domain adaptation using subspace alignment,” in *Proceedings of the IEEE international conference on computer vision*, 2013, pp. 2960–2967.
- [38] J. Cha, W. Kang, J. Mun, and B. Roh, “Honeybee: Locality-enhanced projector for multimodal llm,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 13 817–13 827.
- [39] Y. Zhai, S. Tong, X. Li, M. Cai, Q. Qu, Y. J. Lee, and Y. Ma, “Investigating the catastrophic forgetting in multimodal large language models,” *arXiv preprint arXiv:2309.10313*, 2023.
- [40] B. Liu, Y. Zhu, C. Gao, Y. Feng, Q. Liu, Y. Zhu, and P. Stone, “Libero: Benchmarking knowledge transfer for lifelong robot learning,” *Advances in Neural Information Processing Systems*, vol. 36, pp. 44 776–44 791, 2023.